

CASE STUDY 4:

Code ▾

This study allows us to revisit/renew

1. Regression modeling
2. Properties of Least Squares/Fitting “a line”
3. Multiple observation

Datasets for this study are

1. The main file: gauge.txt
2. Supplementary large-scale files: download the following folder Full Resolution Data.zip More information about the supplementary file can be found at <http://iabp.apl.washington.edu/data.html> (<http://iabp.apl.washington.edu/data.html>) as well as <http://nsidc.org/data/G00791> (<http://nsidc.org/data/G00791>)

Question

The aim of this lab is to provide a simple procedure for converting gain into density when the gauge is in operation. Keep in mind that the experiment was conducted by varying density and measuring the response in gain, but when the gauge is ultimately in use, the snow-pack density is to be estimated from the measured gain.

Setup

Hide

```
df <- read.table('gauge-lwblwa6-2gpel41.txt', header=TRUE)
df <- df[order(df$density), ] # Sort from least to greatest density
m <- 9 # Number of distinct block densities
t <- 10 # Number of replicate measurements
#install.packages('Llpack')
#install.packages('quantreg')
#install.packages('ggplot2')
library(Llpack) # Used for least absolute deviations regression line
library(quantreg) # Used for quantile regression line
library(ggplot2)
```

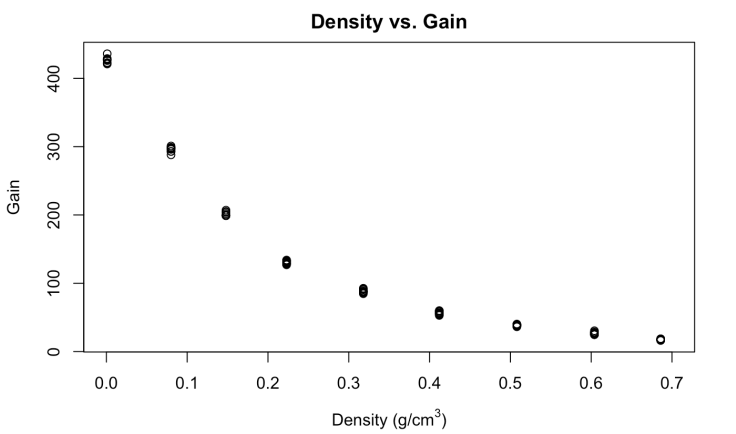
Scenario 1: Fitting

Use the data to fit the gain, or a transformation of gain, to density. Try sketching the least squares line on a scatter plot.

- Do the residuals indicate any problems with the fit?
- If the densities of the polyethylene blocks are not reported exactly, how might this affect the fit?
- What if the blocks of polyethylene were not measured in random order (location)?


Hide

```
# Plot raw data
title <- 'Density vs. Gain'
x.axis <- expression('Density (g/cm3*)')
y.axis <- 'Gain'
x.range <- c(0, .7)
y.range <- c(2.5, 6)
plot(df, main=title, xlab=x.axis, ylab=y.axis, xlim=x.range)
```




Hide

```
# Take log transformation of response variable (gain)
y.log.axis = 'log(Gain)'
df.log = data.frame(df['density'], log(df['gain']))
plot(df.log, main=title, xlab=x.axis, ylab=y.log.axis, xlim=x.range, ylim=y.range)
```



Hide

```
# Average replicate measurements
df.log.avg = aggregate(list(gain=df.log$gain), by=list(density=df.log$density), FUN=mean)
plot(df.log.avg, main=title, xlab=x.axis, ylab=y.log.axis, xlim=x.range, ylim=y.range)
```



Hide

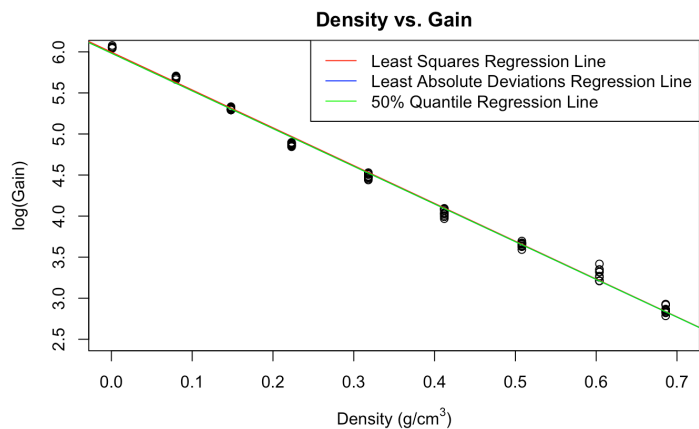
```
# Fit gain to density
least.squares <- lm(gain~density, data=df.log.avg)
lad <- lad(gain~density, data=df.log.avg)
quant <- rq(gain~density, tau=.5, data=df.log.avg)
plot(df.log, main=title, xlab=x.axis, ylab=y.log.axis, xlim=x.range, ylim=y.range)
abline(least.squares, col='red')
```

Hide

```
abline(lad, col='blue')
abline(quant, col='green')
```

Hide

```
legend('topright', legend=c('Least Squares Regression Line', 'Least Absolute Deviations Regression Line', '50% Quantile Regression Line'), col=c('red', 'blue', 'green'), lty=1)
```



Hide

```
c(cor(df.log.avg), summary(least.squares)$r.squared)
```

```
[1] 1.0000000 -0.9984469 -0.9984469 1.0000000 0.9968963
```

Hide

```
least.squares
```

```
Call:
lm(formula = gain ~ density, data = df.log.avg)
```

```
Coefficients:
(Intercept)      density
      5.997        -4.606
```

Hide

```
lad
```

```
Call:
lad(formula = gain ~ density, data = df.log.avg)
Converged in 4 iterations
```

```
Coefficients:
(Intercept)      density
      5.9850        -4.5935
```

```
Degrees of freedom: 9 total; 7 residual
Scale estimate: 0.06926379
```

Hide

```
quant
```

```
Call:
rq(formula = gain ~ density, tau = 0.5, data = df.log.avg)
```

```
Coefficients:
(Intercept)      density
      5.985029        -4.593460
```

```
Degrees of freedom: 9 total; 7 residual
```

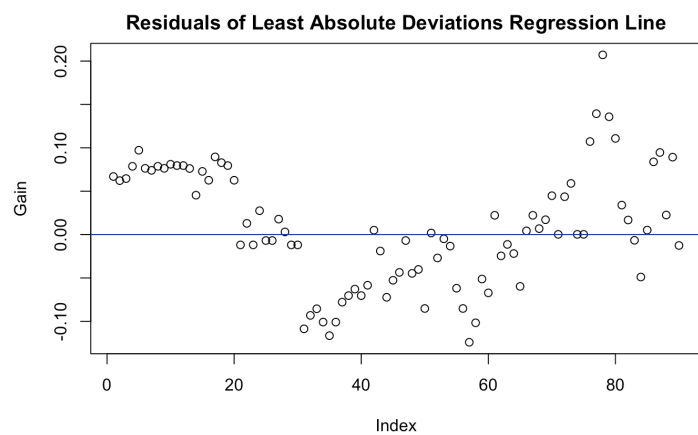
Hide

```
# Check conditions for linear regression: linearity, normality of residuals, and constant variability
least.squares.residuals <- data.frame(df.log['density'], df.log['gain'] - rep(predict(least.squares), each=10))
lad.residuals <- data.frame(df.log['density'], df.log['gain'] - rep(predict(lad), each=10))
quant.residuals <- data.frame(df.log['density'], df.log['gain'] - rep(predict(quant), each=10))
title.residuals1 <- 'Residuals of Least Squares Regression Line'
title.residuals2 <- 'Residuals of Least Absolute Deviations Regression Line'
title.residuals3 <- 'Residuals of 50% Quantile Regression Line'
plot(least.squares.residuals$gain, main=title.residuals1, ylab=y.axis)
abline(0, 0, col='red')
```




Hide

```
plot(lad.residuals$gain, main=title.residuals2, ylab=y.axis)  
abline(0, 0, col='blue')
```



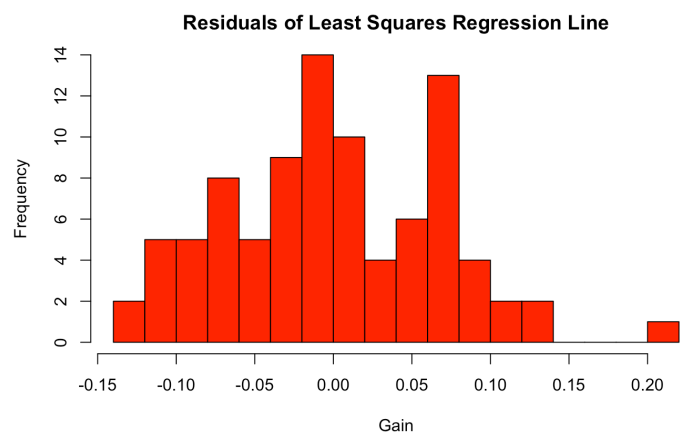
Hide

```
plot(quant.residuals$gain, main=title.residuals3, ylab=y.axis)  
abline(0, 0, col='green')
```



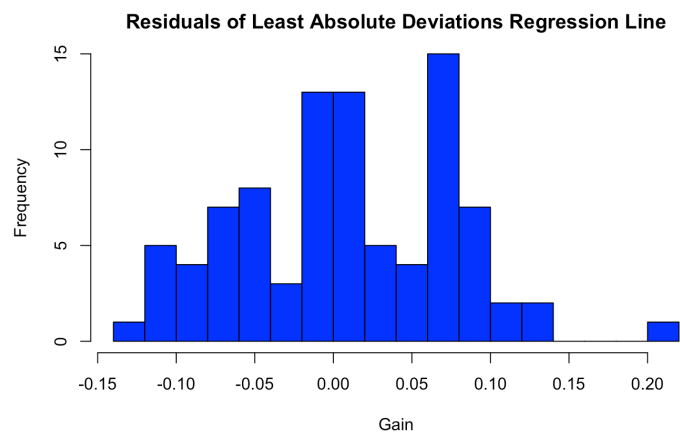
Hide

```
num.bins <- 12  
hist(least.squares.residuals$gain, breaks=num.bins, main=title.residuals1, xlab=y.axis,  
     col='red')
```



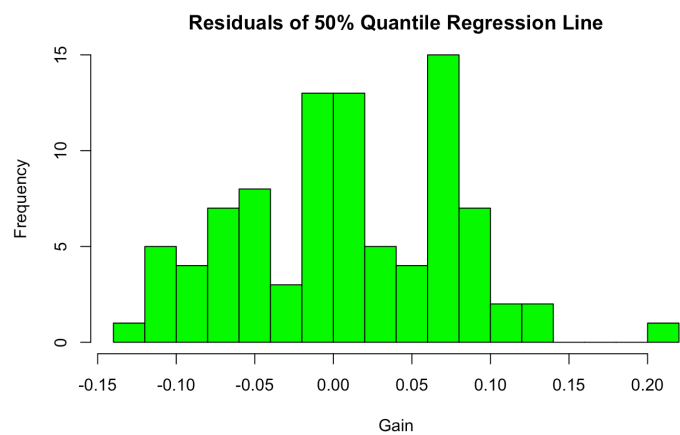
Hide

```
hist(lad.residuals$gain, breaks=num.bins, main=title.residuals2, xlab=y.axis, col='blue')
```




Hide

```
hist(quant.residuals$gain, breaks=num.bins, main=title.residuals3, xlab=y.axis, col='green')
```




Hide

```
qqnorm(least.squares.residuals$gain, main=paste('Normal Q-Q Plot with', title.residuals1), cex.main=1)
qqline(least.squares.residuals$gain, col='red')
```




Hide

```
qqnorm(lad.residuals$gain, main=paste('Normal Q-Q Plot with', title.residuals2), cex.
main=1)
qqline(lad.residuals$gain, col='blue')
```



Hide

```
qqnorm(quant.residuals$gain, main=paste('Normal Q-Q Plot with', title.residuals3), ce
x.main=1)
qqline(quant.residuals$gain, col='green')
```



Scenario 2: Predicting

Ultimately we are interested in answering questions such as: Given a gain reading of 38.6, what is the density of the snow-pack? Given a gain reading of 426.7, what is the density of the snow-pack? These two numeric values, 38.6 and 426.7, were chosen because they are the average gains for the 0.508 and 0.001 densities, respectively.

- Develop a procedure for adding bands around your least squares line that can be used to make interval estimates for the snow-pack density from gain measurements. Keep in mind how the data were collected: several measurements of gain were taken for polyethylene blocks of known density.

Hide

```
# Predictions
PredictLogGain <- function(density)
  predict(least.squares, data.frame(density=density)) # Predict log(gain) using density
PredictDensityLeastSquares <- function(gain) {
  intercept <- coef(least.squares)[[1]]
  slope <- coef(least.squares)[[2]]
  (log(gain) - intercept) / slope # Predict density using gain
}
PredictDensityLad <- function(gain) {
  intercept <- coef(lad)[[1]]
  slope <- coef(lad)[[2]]

  (log(gain) - intercept) / slope # Predict density using gain
}
PredictDensityQuant <- function(gain) {
  intercept <- coef(quant)[[1]]
  slope <- coef(quant)[[2]]
  (log(gain) - intercept) / slope # Predict density using gain
}
# 95% prediction and confidence intervals of log(gain) using density
t <- qt(.975, df=m-2)
mean.density <- mean(df.log.avg$density)
summation <- sum((df.log.avg$density - mean.density) ^ 2)
s2 <- aggregate(list(variance=least.squares$residuals$gain), by=list(density=least.squares$residuals$density), FUN=var)
s.pooled <- sqrt(mean(s2$variance))
center.expr <- quote(center <- PredictLogGain(density))
ci.width.expr <- quote(width <- t * s.pooled * sqrt(1/m + (density-mean.density)^2 / summation))
pi.width.expr <- quote(width <- t * s.pooled * sqrt(1 + 1/m + (density-mean.density)^2 / summation))
LogGainCiLower <- function(density) {
  eval(center.expr)
  eval(ci.width.expr)
  center - width
}
LogGainCiUpper <- function(density) {
  eval(center.expr)
  eval(ci.width.expr)
  center + width
}
LogGainPiLower <- function(density) {
  eval(center.expr)
  eval(pi.width.expr)
  center - width
}
LogGainPiUpper <- function(density) {
  eval(center.expr)
  eval(pi.width.expr)
  center + width
}
# Add bands around least squares line
plot(df.log, main=title, xlab=x.axis, ylab=y.log.axis, xlim=x.range, ylim=y.range)
abline(least.squares, col='red')
```

Hide


```
ci.col <- 'purple'
pi.col <- 'blue'
symbol <- '.'
size <- 1.5
line.type <- 3
line.width <- 0.7
confidence.intervals <- data.frame(density=df.log.avg$density, lower=LogGainCiLower(d
f.log.avg$density), upper=LogGainCiUpper(df.log.avg$density))
points(x=confidence.intervals$density, y=confidence.intervals$lower, col=ci.col, pch=
symbol, cex=size)
points(x=confidence.intervals$density, y=confidence.intervals$upper, col=ci.col, pch=
symbol, cex=size)
```

Hide

```
lines(x=confidence.intervals$density, y=confidence.intervals$lower, col=ci.col, lty=1
ine.type, lwd=line.width)
lines(x=confidence.intervals$density, y=confidence.intervals$upper, col=ci.col, lty=1
ine.type, lwd=line.width)
```

Hide


```
prediction.intervals <- data.frame(density=df.log.avg$density, lower=LogGainPiLower(d
f.log.avg$density), upper=LogGainPiUpper(df.log.avg$density))
points(x=prediction.intervals$density, y=prediction.intervals$lower, col=pi.col, pch=
symbol, cex=size)
points(x=prediction.intervals$density, y=prediction.intervals$upper, col=pi.col, pch=
symbol, cex=size)
```


Hide

```
lines(x=prediction.intervals$density, y=prediction.intervals$lower, col=pi.col, lty=1
ine.type, lwd=line.width)
lines(x=prediction.intervals$density, y=prediction.intervals$upper, col=pi.col, lty=1
ine.type, lwd=line.width)
```

Hide

```
legend('topright', legend=c('Least Squares Regression Line', '95% Confidence Interval
Bands', '95% Prediction Interval Bands'), col=c('red', ci.col, pi.col), lty=1)
```





Hide

```
# 95% prediction and confidence intervals of density using gain
end.points <- c(-1, 3) # Interval to search the root in
DensityCi <- function(gain) {
  lower <- uniroot(function(density) log(gain) - LogGainCiLower(density), interval=en
d.points)[[1]]
  upper <- uniroot(function(density) log(gain) - LogGainCiUpper(density), interval=en
d.points)[[1]]
  c(lower, upper)
}
DensityPi <- function(gain) {
  lower <- uniroot(function(density) log(gain) - LogGainPiLower(density), end.points)
[[1]]
  upper <- uniroot(function(density) log(gain) - LogGainPiUpper(density), end.points)
[[1]]
  c(lower, upper)
}
# Point and interval estimates for example gain readings
PredictDensityLeastSquares(38.6) # 38.6 is the average gain for 0.508 density
```

[1] 0.5089113

Hide

```
PredictDensityLad(38.6)
```

[1] 0.5076298

Hide

```
PredictDensityQuant(38.6)
```

[1] 0.5076298

Hide

```
DensityCi(38.6)
```

[1] 0.5011879 0.5169000

Hide

```
DensityPi(38.6)
```

[1] 0.4889568 0.5291323

Hide

```
PredictDensityLeastSquares(426.7) # 426.7 is the average gain for 0.001 density
```

[1] -0.01276954

Hide

```
PredictDensityLad(426.7)
```

[1] -0.01546807

Hide

```
PredictDensityQuant(426.7)
```

[1] -0.01546811

Hide

```
DensityCi(426.7)
```

[1] -0.024285866 -0.001769193

Hide

```
DensityPi(426.7)
```

```
[1] -0.034644666 0.008618629
```

Scenario 3: Cross-Validation

To check how well your procedure works, omit the set of measurements corresponding to the block of density 0.508, apply your “estimation”/calibration procedure to the remaining data, and provide an interval estimate for the density of a block with an average reading of 38.6. Where does the actual density fall in the interval? Try the same test, for the set of measurements at the 0.001 density.

Hide

```
for (omitted in c(0.508, 0.001)) {
  # Omit measurements corresponding to the specified density
  df.log.omitted = df.log[which(df.log['density'] != omitted), ]
  df.log.avg.omitted <- df.log.avg[which(df.log.avg['density'] != omitted), ]

  # Redo calculations using modified dataset
  least.squares <- lm(gain~density, data=df.log.avg.omitted)

  mean.density <- mean(df.log.avg.omitted$density)
  summation <- sum((df.log.avg.omitted$density - mean.density) ^ 2)

  s2 <- aggregate(list(variance=least.squares$residuals$gain, by=list(density=least.squares$residuals$density), FUN=var)
  s.pooled <- sqrt(mean(s2$variance))

  ci.width.expr <- quote(width <- t * s.pooled * sqrt(1/(m-1) + (density-mean.density)^2 / summation))
  pi.width.expr <- quote(width <- t * s.pooled * sqrt(1 + 1/(m-1) + (density-mean.density)^2 / summation))

  plot(df.log.omitted, main=title, xlab=x.axis, ylab=y.log.axis, xlim=x.range, ylim=y.range)
  abline(least.squares, col='red')

  ci.col <- 'purple'
  pi.col <- 'blue'
  symbol <- '-'
  size <- 1.5
  line.type <- 3
  line.width <- 0.7

  confidence.intervals <- data.frame(density=df.log.avg.omitted$density, lower=LogGainCiLower(df.log.avg.omitted$density), upper=LogGainCiUpper(df.log.avg.omitted$density))
  points(x=confidence.intervals$density, y=confidence.intervals$lower, col=ci.col, pch=symbol, cex=size)
  points(x=confidence.intervals$density, y=confidence.intervals$upper, col=ci.col, pch=symbol, cex=size)
  lines(x=confidence.intervals$density, y=confidence.intervals$lower, col=ci.col, lty=line.type, lwd=line.width)
  lines(x=confidence.intervals$density, y=confidence.intervals$upper, col=ci.col, lty=line.type, lwd=line.width)

  prediction.intervals <- data.frame(density=df.log.avg.omitted$density, lower=LogGainPiLower(df.log.avg.omitted$density), upper=LogGainPiUpper(df.log.avg.omitted$density))
  points(x=prediction.intervals$density, y=prediction.intervals$lower, col=pi.col, pch=symbol, cex=size)
  points(x=prediction.intervals$density, y=prediction.intervals$upper, col=pi.col, pch=symbol, cex=size)
  lines(x=prediction.intervals$density, y=prediction.intervals$lower, col=pi.col, lty=line.type, lwd=line.width)
  lines(x=prediction.intervals$density, y=prediction.intervals$upper, col=pi.col, lty=line.type, lwd=line.width)


  legend('topright', legend=c('Least Squares Regression Line', '95% Confidence Interval Bands', '95% Prediction Interval Bands'), col=c('red', ci.col, pi.col), lty=1)


  print(PredictDensityLeastSquares(38.6)) # 38.6 is the average gain for 0.508 density
  print(DensityCi(38.6))
  print(DensityPi(38.6))

  print(PredictDensityLeastSquares(426.7)) # 426.7 is the average gain for 0.001 density
  print(DensityCi(426.7))
  print(DensityPi(426.7))
}
```

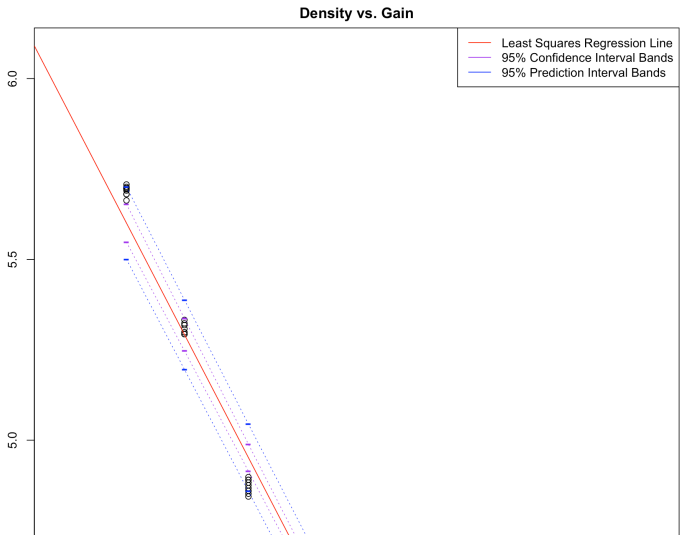
```
[1] 0.5091927
[1] 0.5006695 0.5180406
[1] 0.4889184 0.5297925
[1] -0.0128045
[1] -0.024342164 -0.001790802
[1] -0.034760736 0.008598777
```


Density vs. Gain





```
[1] 0.5092919
[1] 0.5014370 0.5174323
[1] 0.4890257 0.5298473
[1] -0.02051733
[1] -0.035344991 -0.006521825
[1] -0.044604850 0.002738127
```





Additional Scenario: Temperature, DOY, and Latitude.

Use the additional dataset to construct a model fitting temperature with DOY, latitude, and other reasonable features. Try sketching the least squares line on a scatter plot. We aim to investigate the relationship between temperature and the DOY, and its latitude.

Hide

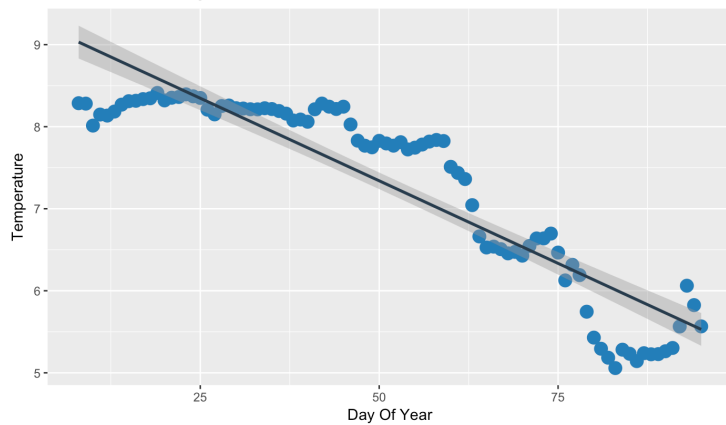
```
# Check the correlation
data <- read.csv('Full Resolution Data/64506420.csv', header=TRUE)
data <- data[,c('Hour','DOY','POS_DOY','Lat','Lon','Ts','BP')]
# Drop the extreme outlier case
#data <- data[which(data$Ts>-200),]
data_matrix <- as.matrix(data)
# Correlation Matrix
corr_matrix <- cor(data_matrix)
corr_matrix
```

	Hour	DOY	POS_DOY	Lat	Lon	Ts
BP						
Hour	1.000000000	-0.006779489	-0.006775002	-0.007511024	0.01217860	0.008592576
0.02505819						
DOY	-0.006779489	1.000000000	0.999999958	0.903704617	-0.68012490	-0.906918658
-0.17232397						
POS_DOY	-0.006775002	0.999999958	1.000000000	0.903696909	-0.68013445	-0.906916964
-0.17230481						
Lat	-0.007511024	0.903704617	0.903696909	1.000000000	-0.59748962	-0.958835395
-0.29666821						
Lon	0.012178596	-0.680124899	-0.680134450	-0.597489620	1.00000000	0.564136723
-0.02981279						
Ts	0.008592576	-0.906918658	-0.906916964	-0.958835395	0.56413672	1.000000000
0.20223136						
BP	0.025058188	-0.172323969	-0.172304805	-0.296668215	-0.02981279	0.202231363
1.000000000						

Hide

```
# Group by DOY and average replicated measurements
data$DOY <- round(data$DOY,0)
data.avg <- aggregate(list(data=data[,c('Ts','Lat')])), by=list(DOY=data$DOY), FUN=mean
)
# least squares line
ggplot(data.avg,aes(x=data.avg$DOY, y=data.avg$data.Ts)) +
  geom_point(color='#2980B9', size = 4) +
  geom_smooth(method=lm, color='#2C3E50') +ggtitle(label ="Least Squares Regression L
ine") + xlab("Day Of Year") +
  ylab("Temperature")
```

Least Squares Regression Line



Hide

```
fit1<-lm(formula = data.Ts ~ DOY, data = data.avg)
summary(fit1)
```

Call:
lm(formula = data.Ts ~ DOY, data = data.avg)

Residuals:

Min	1Q	Median	3Q	Max
-0.95439	-0.34633	0.03139	0.35247	0.84693

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.353218	0.114242	81.87	<2e-16 ***
DOY	-0.040253	0.001989	-20.23	<2e-16 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4741 on 86 degrees of freedom
Multiple R-squared: 0.8264, Adjusted R-squared: 0.8244
F-statistic: 409.4 on 1 and 86 DF, p-value: < 2.2e-16

Hide

```
ggplot(data.avg,aes(x=data.avg$data.Lat, y=data.avg$data.Ts)) +
  geom_point(color='#2980B9', size = 4) +
  geom_smooth(method=lm, color='#2C3E50') +ggtitle(label ="Least Squares Regression L
ine")+ xlab("Latitude") +
  ylab("Temperature")
```

Least Squares Regression Line



Hide

```
fit2<-lm(formula = data.Ts ~ data.Lat, data = data.avg)
summary(fit2)
```

```
Call:
lm(formula = data.Ts ~ data.Lat, data = data.avg)

Residuals:
    Min       1Q   Median       3Q      Max
-0.51377 -0.26924 -0.04201  0.24403  0.64902

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  47.9901     1.2672   37.87  <2e-16 ***
data.Lat     -0.6202     0.0193  -32.14  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3155 on 86 degrees of freedom
Multiple R-squared:  0.9231,    Adjusted R-squared:  0.9222
F-statistic: 1033 on 1 and 86 DF,  p-value: < 2.2e-16
```

Hide

```
# Polynomial Regression Line
fit3<-lm(formula = data.Ts ~ DOY + data.Lat, data = data.avg)
summary(fit3)
```

```
Call:
lm(formula = data.Ts ~ DOY + data.Lat, data = data.avg)


Residuals:
    Min       1Q   Median       3Q      Max
-0.60721 -0.22496 -0.00537  0.25822  0.61356

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  40.049638     2.673979   14.978  < 2e-16 ***
DOY          -0.009756     0.002936   -3.322  0.00132 **
data.Lat     -0.491566     0.042805  -11.484  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2985 on 85 degrees of freedom
Multiple R-squared:  0.932, Adjusted R-squared:  0.9304
F-statistic: 582.1 on 2 and 85 DF,  p-value: < 2.2e-16
```

Hide

```
qqnorm(fit2$residuals, main=paste('Normal Q-Q Plot with', title.residuals1), cex.main
=1)
qqline(fit2$residuals, col='red')
```



Hide

```
title.residuals1 <- 'Residuals of Least Square Regression Line'
plot(fit2$residuals, main=title.residuals1, ylab = "Standardized Residuals")
abline(0, 0, col='red')
```

