Tree, Bagging, Random Forests and Boosting

Only used for predictions

not CI / Testing / significant variables.

- Classification trees ← Leo Breiman
- Boosted Random Forests ← run a lot of trees and find one that is not in the forest but better than all existing trees in the forests

Non-parametric

Arm Bended Problem

Two-class Classifications

If tree is small ⇒ good prediction properties.

⋮ ⇒ how to determine when ⇒ in R package
to stop

huge ⇒ overfit

SPAM ⇒ R built-in

Sensitivity ⇒ proportion of true spam identified ↑ Type I ↑

Specificity ⇒ proportion of true email identified ↑ Power

Want both to be high

overfitting ⇒ could get high specificity & sensitivity

Freund → prof here

Decision Boundary : Tree

Model Averaging ⇒ Boosting > Random Forests > Bagging > Single Tree

Bagging

Bootstrap aggregation

⇒ bootstrap a thousand times

⇒ get a thousand trees

⇒ average the trees

Smoother decision boundaries

Random Forests:

⇒ Randomly choose m features (refinement)

then ↳ bagging ⇒ average the trees

⇒ high dimensional

Boosting

⇒ Bootstrap, take features

⇒ which data was predicted well ⇒ more likely ⇒ smaller weight

⇒ higher weight on misfit ⇒ weighted bootstrap ⇒ run trees.

⇒ which data predict well ⇒ reweight

↳ Loss function

pick the data you misclassified

⇒ won't be overfitting

run many forests

⇒ converged forests